Computer Vision - Simpsons Family - Character Recognition

ANDREI-CRISTIAN RUSU

January 7, 2022

1 Introduction

In this document you will find implementation details about the Simpsons Family Character Recognition application.

2 Reading the training data

For training, we have the following directory structure:

```
data_task1
  └── [positive/negative] // type of image
      └── [character name]_[index].jpg // image
data_task2
  └── [bart/homer/lisa/marge/unknown] // character name
      └── [character name]_[index].jpg // image
```

The fact that we're duplicating the images is not the optimal way to do things. The only reason we do this is to simplify things with the *PyTorch data loader*. This will be updated in the future.

3 Generating the data

The generators for positive / negative images can be found in *ExamplesGenerator.py*. It contains two functions:

- *generate_positive_examples()*

- *generate_negative_examples(count_per_image)*

The first function:

- iterates through every image in the train directory

- finds the corresponding box for each detection

- crops that box

- resizes it to $(image\_height, image\_width)$ (i.e. to $(72, 72)$)

- saves it to the corresponding directory

The second one:

- iterates through every image in the train directory

- finds the corresponding boxes for each detection

- for *count_per_image* times, it tries to crop random patches which do not intersect with the real detections *or* have at most 50% yellow pixels.

- for the other third of *count_per_image*, it crops random patches that might or might not intersect with the real detection, but that contain at least 1% and at most 3% yellow pixels

- saves each patch to the corresponding directory

4 Training the Neural Networks

We mainly have two tasks:

- Given a image of $n * m$ size, check if there is a face of a character.

- Given a image of $n * m$ size which contains a character, determine the name of the character.

In order to solve this, we will train two neural networks:

The first CNN will be trained on the data from *data_task*1 directory - and will learn how a *positive* image looks like (an image which contains a character).

The second CNN will be trained on data from *data_task*2 - learning the name of the characters.

For both of our networks, we'll be using *resnet18*, with a *SGD* optimizer and *Early Stopping ReduceLROnPlateau*. The implementation resides in *NeuralNetwork.py*.

The hyperparameters that we chose are the following:

- batch_size: 64

- momentum: 0.9

- learning_rate: $1e - 2$

- early_stopping_patience: 3

- reduce_lr_on_plateau_patience: 1

The number of epochs is *at most* 50.

When training the network, we are *augmenting* the images, using the following transformations:

```
1   transform: Compose = Compose([
2       Resize((image_height, image_width)),
3
4       RandomHorizontalFlip(),
5       RandomRotation(degrees=10),
6       RandomAffine(degrees=0.2),
7
8       ToTensor(),
9   ])
```

Here are some logs from the training session:

```
1   2022-01-07 07:38:46,401/INFO: Training CNN #1...
2   2022-01-07 07:39:29,345/INFO: Epoch 1 | loss: 0.09360 | acc: 0.96752 | val_loss: 0.05067 | val_acc: 0.98376
3   2022-01-07 07:40:10,040/INFO: Epoch 2 | loss: 0.04641 | acc: 0.98489 | val_loss: 0.03776 | val_acc: 0.98790
4   2022-01-07 07:40:51,098/INFO: Epoch 3 | loss: 0.04118 | acc: 0.98641 | val_loss: 0.03156 | val_acc: 0.98986
```

5 2022-01-07 07:41:32,412/INFO: Epoch 4 | loss: 0.03114 | acc: 0.98997 | val_loss: 0.03196 | val_acc: 0.98923
6 EarlyStopping counter: 1 out of 3
7 2022-01-07 07:41:32,476/INFO: Done.
8 2022-01-07 07:41:32,476/INFO: Training CNN #2...
9 2022-01-07 07:41:48,172/INFO: Epoch 1 | loss: 1.05340 | acc: 0.59094 | val_loss: 0.55163 | val_acc: 0.81775
10 2022-01-07 07:42:03,595/INFO: Epoch 2 | loss: 0.42271 | acc: 0.85790 | val_loss: 0.33255 | val_acc: 0.89311
11 2022-01-07 07:42:18,857/INFO: Epoch 3 | loss: 0.28836 | acc: 0.90759 | val_loss: 0.23276 | val_acc: 0.92629
12 2022-01-07 07:42:34,724/INFO: Epoch 4 | loss: 0.22213 | acc: 0.92978 | val_loss: 0.17778 | val_acc: 0.94811
13 2022-01-07 07:42:50,113/INFO: Epoch 5 | loss: 0.17843 | acc: 0.94591 | val_loss: 0.14874 | val_acc: 0.95728
14 2022-01-07 07:43:05,548/INFO: Epoch 6 | loss: 0.16297 | acc: 0.95160 | val_loss: 0.11749 | val_acc: 0.96700
15 2022-01-07 07:43:21,109/INFO: Epoch 7 | loss: 0.11551 | acc: 0.96956 | val_loss: 0.09829 | val_acc: 0.97158
16 2022-01-07 07:43:36,593/INFO: Epoch 8 | loss: 0.09721 | acc: 0.97140 | val_loss: 0.07810 | val_acc: 0.97910
17 2022-01-07 07:43:36,654/INFO: Done.

5 Recognizing faces

For this, we're using the $run_face_detector function from FaceDetector.py$.

The logic behind this function is the following:

- *iterate through every image in the test directory*

- *for every image, create an array of detections and scores*

- *iterate through a number of different $n*n$ window_sizes*

- *for each window, apply the sliding window technique, and get the (x, y) coordinates of the top left corner*

- *if we can't create a $(window\_size, window\_size)$ window, continue*

- *get the percentage of yellow pixels from the window, and, if it's not in $[0.3, 0.7]$, continue*

- *run the detection CNN on the window, and get the score class*

- *if the window is a positive detection, add the coordinates of the window the score to the arrays*

- *after iterating with each window, with all window sizes, we continue if there were no detections. otherwise,*

- *we apply non maximum suppresion over all detections, using the associated scores, and, for each detection,*

- *we run the second CNN on the corresponding window and get the name of the character that is in the image*

- *we add all the needed data to the corresponding arrays, and, at the end, return the filenames of the images, with the detections, scores and names.*

6   End notes

All the hyper parameters that were used in the application were found either by manually testing, or via a Grid Search technique.

   Everything can be found in the implementation.